


0x13 - Memory and Caches

ENGR 3410: Computer Architecture


Jon Tse

Fall 2020


Feedback - Workload




Feedback - Understanding



Feedback - Pace




Feedback - Tools



Feedback - More Time...

I wish we spent more time on...

10 responses



Feedback - Top of Mind?

Anything else on your mind?


1 response

I'm having a lot of fun debugging the CPU

Final Project

- NINJAs will reduce hours, but...
- You can ask a NINJA to be your “advisor!”
- Reach out, they want to learn with you!

Review



- So far, we have three types of memory:
 - Registers/DFF (between pipeline stages)
 - Register File
 - Black box memory for loads and stores
- Today we'll investigate the internals of the memory system

Goals

- As always, computer architects care about:
Performance – Area – Energy – *Design Time*
- For memory, you can have
 - Fast but Expensive (large area)
 - Slow but Cheap
- Our goal – design the best of both worlds


Static Random Access Memory

(SRAM)


- Like a register file, but:
 - Only one Port (unified read/write)
 - Many more words (greater depth)
- Different Cell Construction
 - Smaller than a D-Flip Flop
- Different Word Select Construction
 - Individual cells are weak, so we add circuitry at the perimeter to amplify during a read

SRAM Cell


- **M1-M3** store bit weakly
- **Word Line** accesses cell
 - Open M5, M6
- **Bit Lines**
 - Read bit (weak signal)
 - Write bit (strong signal)
- Cross-coupled like a latch



Dynamic DRAM Cell




- Capacitor stores bit
 - For a while
 - Must be “refreshed”
- FET controls access
- Smaller than SRAM
- Slower than SRAM



http://www.emrl.de/imagesArticles/DRAM_Emulation_Fig2.jpg


Flash Cell

- Transistor with floating gate stores bit
 - Nonvolatile!
 - Limited lifespan
- Possible to have multiple bits per cell
- Requires elevated voltage for program/erase



Technology Trends

- Processor-DRAM Memory Gap (latency)



Technology Trends – Modern Story




Figure from Hennessy and Patterson's Computer Architecture: A Quantitative Approach

The Problem

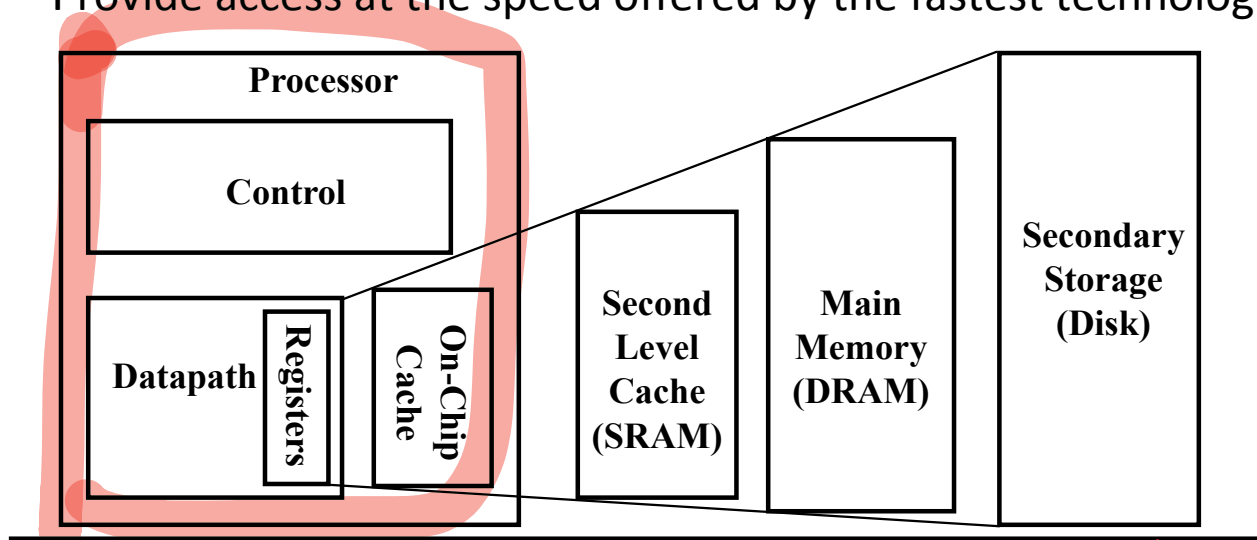
- The Von Neumann Bottleneck *"Memory Wall"*
 - Logic gets faster
 - Memory capacity gets larger
 - **Memory speed is not keeping up with logic**
- How do we cope?

Fast, Big, Cheap: Pick 3

- Design Philosophy
 - Use a hybrid approach that uses aspects of both
 - Lots of slow + cheap
 - Small amount of fast + expensive
- “Cache”
 - Make the common case fast
 - Keep frequently used things in a small amount of fast/expensive memory

The Solution

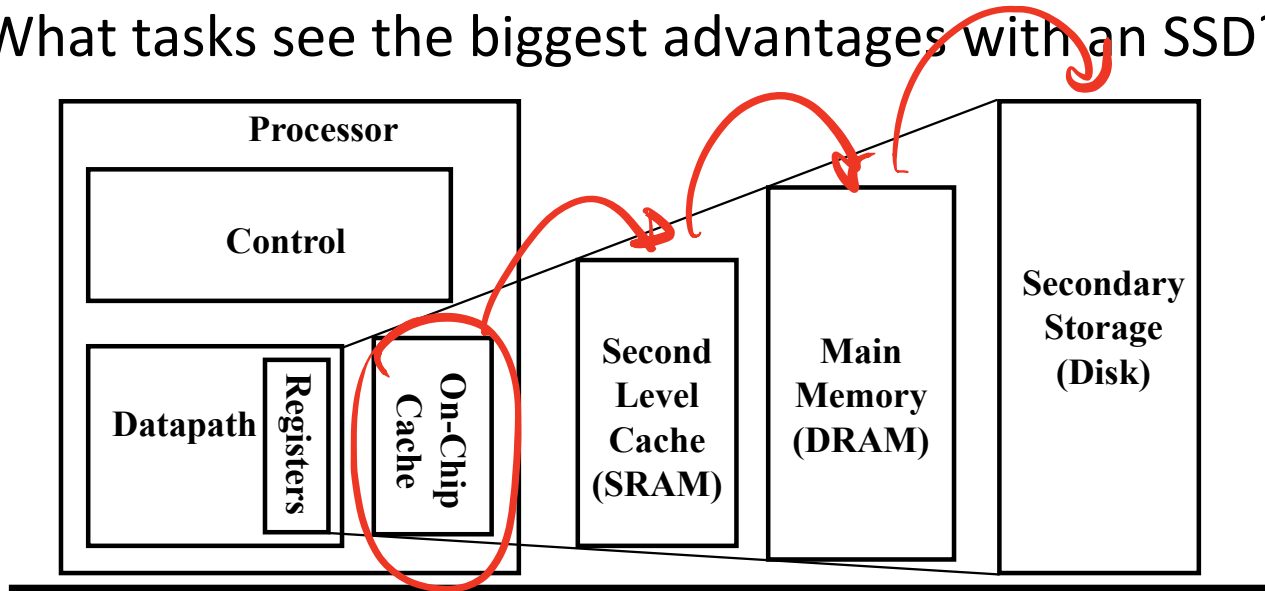
- By taking advantage of the principle of locality:
 - Provide as much memory as is available in the cheapest technology.
 - Provide access at the speed offered by the fastest technology.



Type	<i>DRP</i> Register	<i>SRAM</i> Cache	<i>DRAM</i> Main Memory	<i>magnetic</i> Disk
Speed	< 1ns	< 10ns	<u>~60ns</u>	<u>10 ms</u>
Size	100B	KB-MB	GB	TB

SSD

- 100-1000x faster access time than HDD
- 3-8x faster read/write **rate**
- What tasks see the biggest advantages with an SSD?




Type	Regfile	Cache	Main Memory	SSD
Speed	< 1ns	< 10ns	~60ns	0.1 ms
Size	100B	KB-MB	GB	TB

Cache Terminology

- **Hit:** Data appears in that level
 - **Hit rate** – percent of accesses hitting in that level
 - **Hit time** – Time to access this level
 - Hit time = Access time + Time to determine hit/miss
- **Miss:** Data does not appear in that level and must be fetched from lower level
 - **Miss rate** – percent of misses at that level = $(1 - \text{hit rate})$
 - **Miss penalty** – Overhead in getting data from a lower level
 - Miss penalty = Lower level access time + Replacement time + Time to deliver to processor
- Miss penalty is usually MUCH larger than the hit time

Cache Access Time



- Average access time
 - Access time = (hit time) + (miss penalty) x (miss rate)
 - Want low miss rate & low hit time, since miss penalty is large
- Average Memory Access Time (AMAT)
 - Apply average access time to entire hierarchy.

Split Caches

- Often separate Instruction and Data Caches
 - Harvard Architecture: Split I & D
 - von Neumann Architecture: Unified
- Higher bandwidth
- Optimize to usage
- Slightly higher miss rate: each cache is smaller

Handling A Cache Miss

- Data Miss
 1. Stall the pipeline (freeze following instructions)
 2. Instruct memory to perform a read and wait
 3. Return the result from memory and allow the pipeline to continue
- Instruction Miss
 1. Send the original PC to the memory
 2. Instruct memory to perform a read and wait (no write enables)
 3. Write the result to the appropriate cache line
 4. Restart the instruction

ref
L1 }
L2 }
B }
DRAW

Cache Access Time Example

Level	Hit Time	Hit Rate	Access Time
L1	1 cycle	95%	
L2	10 cycles	90%	
Main Memory	50 cycles	99%	
Disk	50,000 cycles	100%	

- Note: Numbers are **local** hit rates – the ratio of access that go to that cache that hit (remember, higher levels filter accesses to lower levels)

Cache Access Time Example

Level	Hit Time	Hit Rate	Access Time
L1	1 cycle	95%	
L2	10 cycles	90%	
Main Memory	50 cycles	99%	
Disk	50,000 cycles	100%	50,000

- Note: Numbers are **local** hit rates – the ratio of access that go to that cache that hit (remember, higher levels filter accesses to lower levels)

Cache Access Time Example

Level	Hit Time	Hit Rate	Access Time
L1	1 cycle	95%	
L2	10 cycles	90%	
Main Memory	50 cycles	99%	$50 + .01 * 50000 = 550$
Disk	50,000 cycles	100%	50,000

- Note: Numbers are **local** hit rates – the ratio of access that go to that cache that hit (remember, higher levels filter accesses to lower levels)

Cache Access Time Example

Level	Hit Time	Hit Rate	Access Time
L1	1 cycle	95%	
L2	10 cycles	90%	$10 + .1 * 550 = 65$
Main Memory	50 cycles	99%	$50 + .01 * 50000 = 550$
Disk	50,000 cycles	100%	50,000

- Note: Numbers are **local** hit rates – the ratio of access that go to that cache that hit (remember, higher levels filter accesses to lower levels)

Cache Access Time Example

Level	Hit Time	Hit Rate	Access Time
L1	1 cycle	95%	$1 + .05 * 65 = 4.25$
L2	10 cycles	90%	$10 + .1 * 550 = 65$
Main Memory	50 cycles	99%	$50 + .01 * 50000 = 550$
Disk	50,000 cycles	100%	50,000

- Note: Numbers are **local** hit rates – the ratio of access that go to that cache that hit (remember, higher levels filter accesses to lower levels)

How do we get those Hit Rates?

- There's no such thing as a free lunch
 - If access was totally random, we'd be out of luck
 - But we have knowledge a priori about programs
- Locality
 - Temporal Locality – If an item has been accessed recently, it will tend to be accessed again soon
 - Spatial Locality – If an item has been accessed recently, nearby items will tend to be accessed soon

Example

```
char *index = string;
while (*index != 0) { /* C strings end in 0 */
    if (*index >= 'a' && *index <= 'z')
        *index = *index + ('A' - 'a');
    index++;
}
```

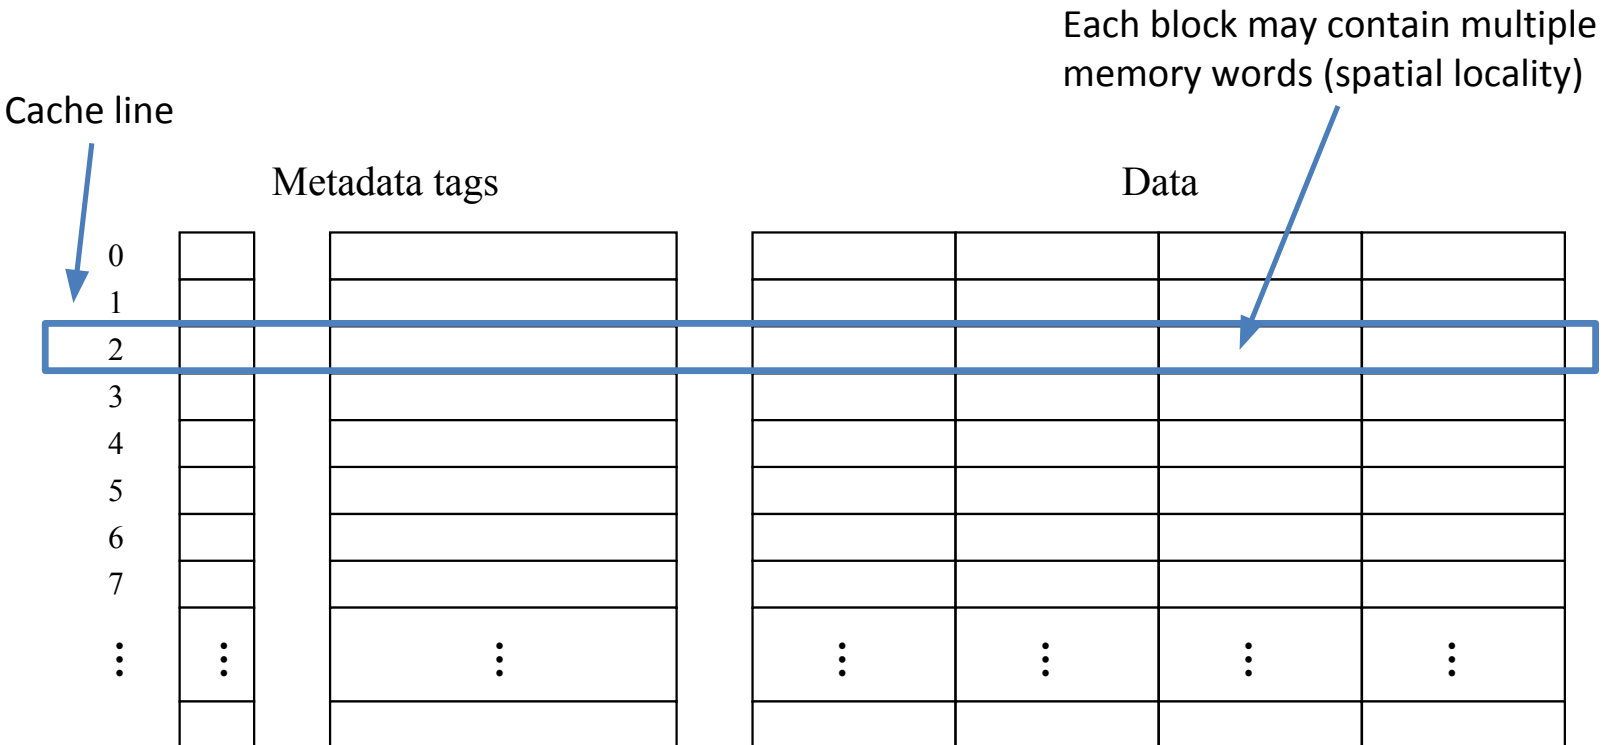
- What does this code do?
- What type(s) of locality does it have?

Exploiting Locality

- Temporal locality
 - Keep more recently accessed items closer to the processor
 - When we must evict items to make room for new ones, attempt to keep more recently accessed items
- Spatial locality
 - When there is a cache miss, move blocks of multiple contiguous words into cache

Whose Line is it, Anyway?

How do we map System Memory into Cache?



Design questions

Assume a fixed amount of storage for your cache design (data + metadata). Propose a cache organization, answering the following questions:

- When moving data from memory into cache, where (i.e. which cache line) do we put it?
- How do we detect whether data for a given memory address is in the cache or not?
- What happens when we run out of cache space?
- How many memory words should go into each cache block?

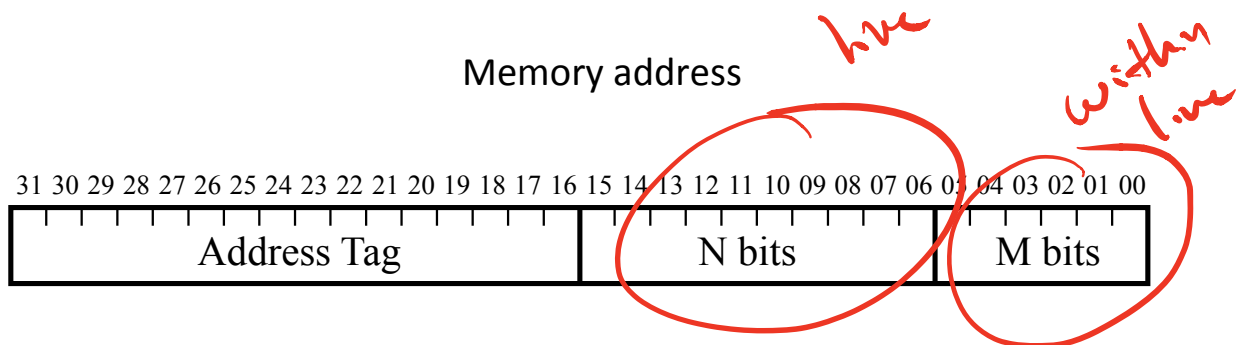
Example Cache Design

- 2^N **Blocks** of Data 2^M bytes wide
- **Tag** indicates what is stored in each block

Valid Bit		Tag (Address)	Data			
0			0	1	...	2^M-1
1						
2						
3						
4						
5						
6						
7						
⋮	⋮	⋮	⋮	⋮	⋮	⋮
2^n-1						

Idea One: Direct Mapping

- Each address maps to a single cache line
- Lower M bits are address within a block
- Next N bits are the cache line address
- Remainder are the address Tag



Direct Mapped Cache

System Memory

0000

0001

0010

0011

0100

0101

0110

0111

1000

1001

1010

1011

1100

1101

1110

1111

Tag

	00
	01
	10
	11

Data

0

1

03 02 01 00

0	00	0
---	----	---

Tag

Line

Block



Direct Mapped Cache

System Memory

0000	
0001	
0010	
0011	
0100	
0101	
0110	
0111	
1000	
1001	
1010	
1011	
1100	
1101	
1110	
1111	


Tag		Data	
	00		
	01		
	10		
	11		
		0	1

03	02	01	00
0	0	0	0
Tag	Line	Block	

Memory address split into Tag, Line, and Block (sizes N,M depend on sizes chosen for cache)

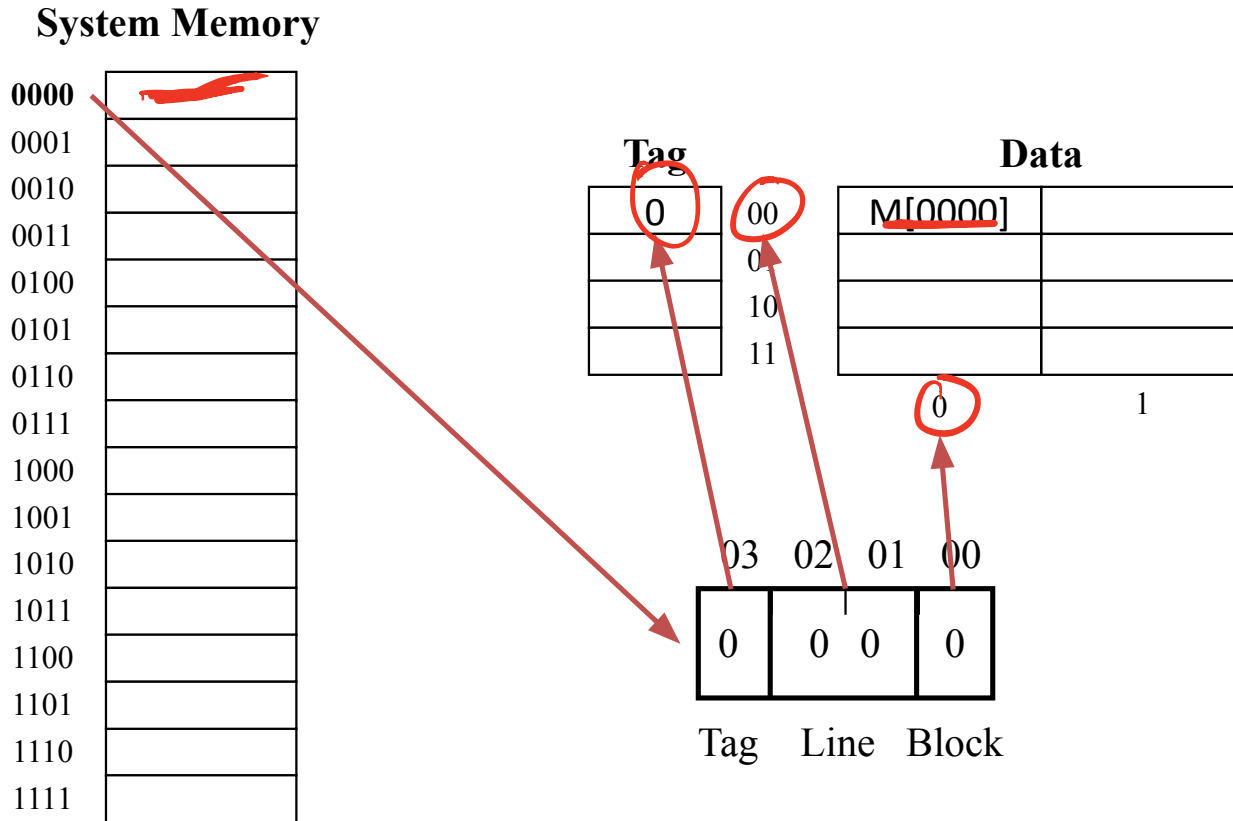
Direct Mapped Cache

System Memory



Line portion of address determines which cache line to use (direct mapped)

Direct Mapped Cache



Requested data loaded from system memory into cache, tag metadata set

Direct Mapped Cache

System Memory

0000	
0001	
0010	
0011	
0100	
0101	
0110	
0111	
1000	
1001	
1010	
1011	
1100	
1101	
1110	
1111	

Tag

0

Data

M[0000]	M[0001]


03	02	01	00
0	0	0	0

Tag Line Block

0 0 0 1

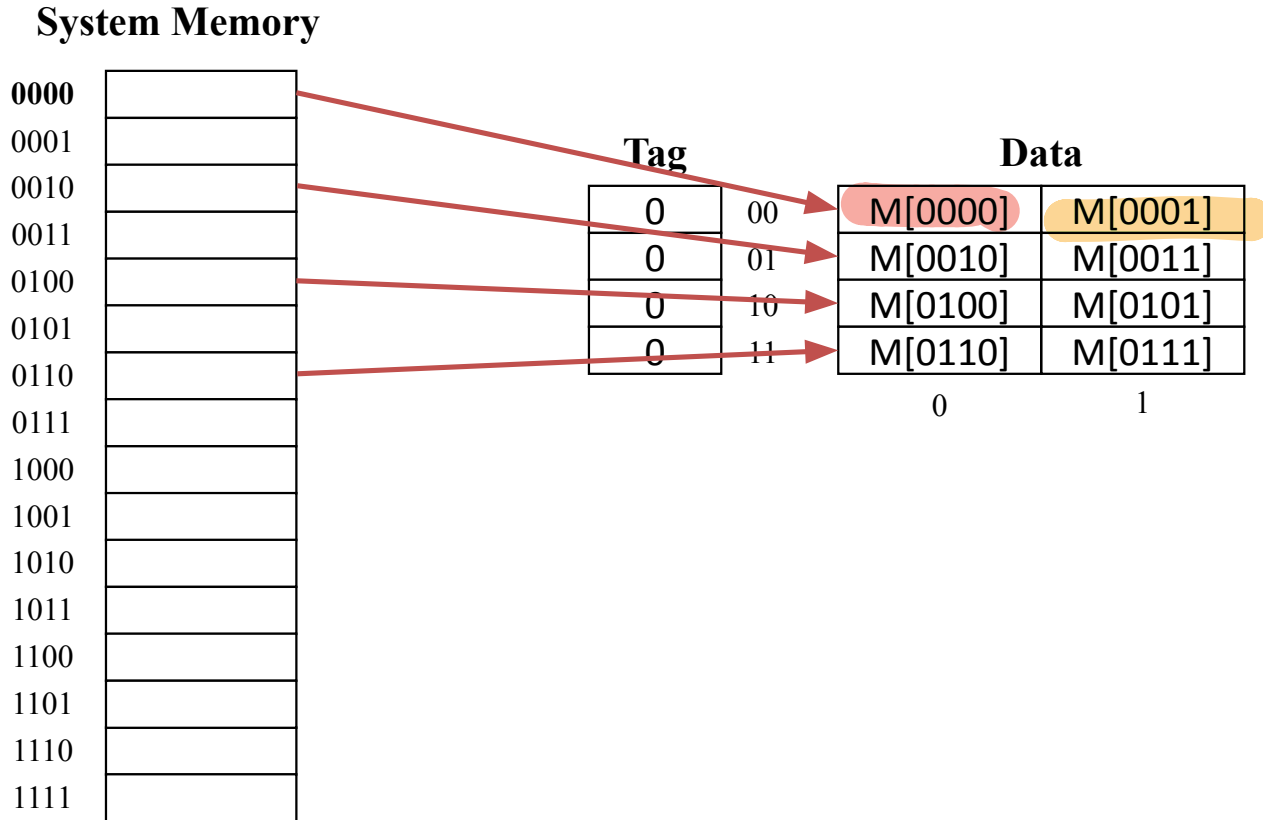
Entire cache block is loaded into cache when any element is read (spatial locality)

Direct Mapped Cache



Only one possible cache line location for each memory address (direct mapping)


Direct Mapped Cache



Cache filled (sequential accesses yield 4 misses, 4 hits)

Direct Mapped Cache

System Memory



Conflict: cache is smaller than memory, so these addresses map to the same cache lines

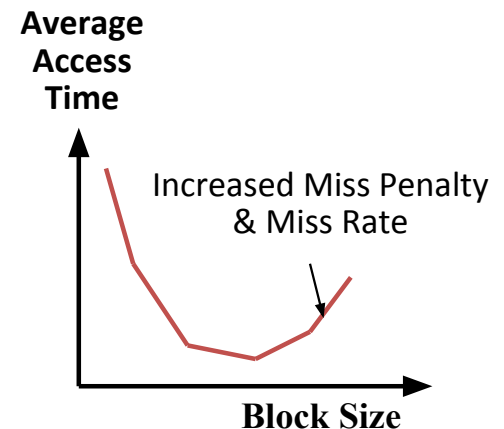


Direct Mapping

- Simple Control Logic
 - Great Spatial Locality
 - Awful Temporal Locality
-
- When does this fail in a direct mapped cache?

```
char *image1, *image2;
int stride;
for(int i=0;i<size;i+=stride) {
    diff +=abs(image1[i] - image2[i]);
}
```


Block Size Tradeoff

- With **fixed cache size**, increasing Block size(M):
 - Worse Miss penalty: longer to fill block
 - Can be partially be hidden by sending the requested data first
 - Better Spatial Locality
 - Worse Temporal Locality



N-Way Set Associative Mapping

- Like N parallel direct map caches

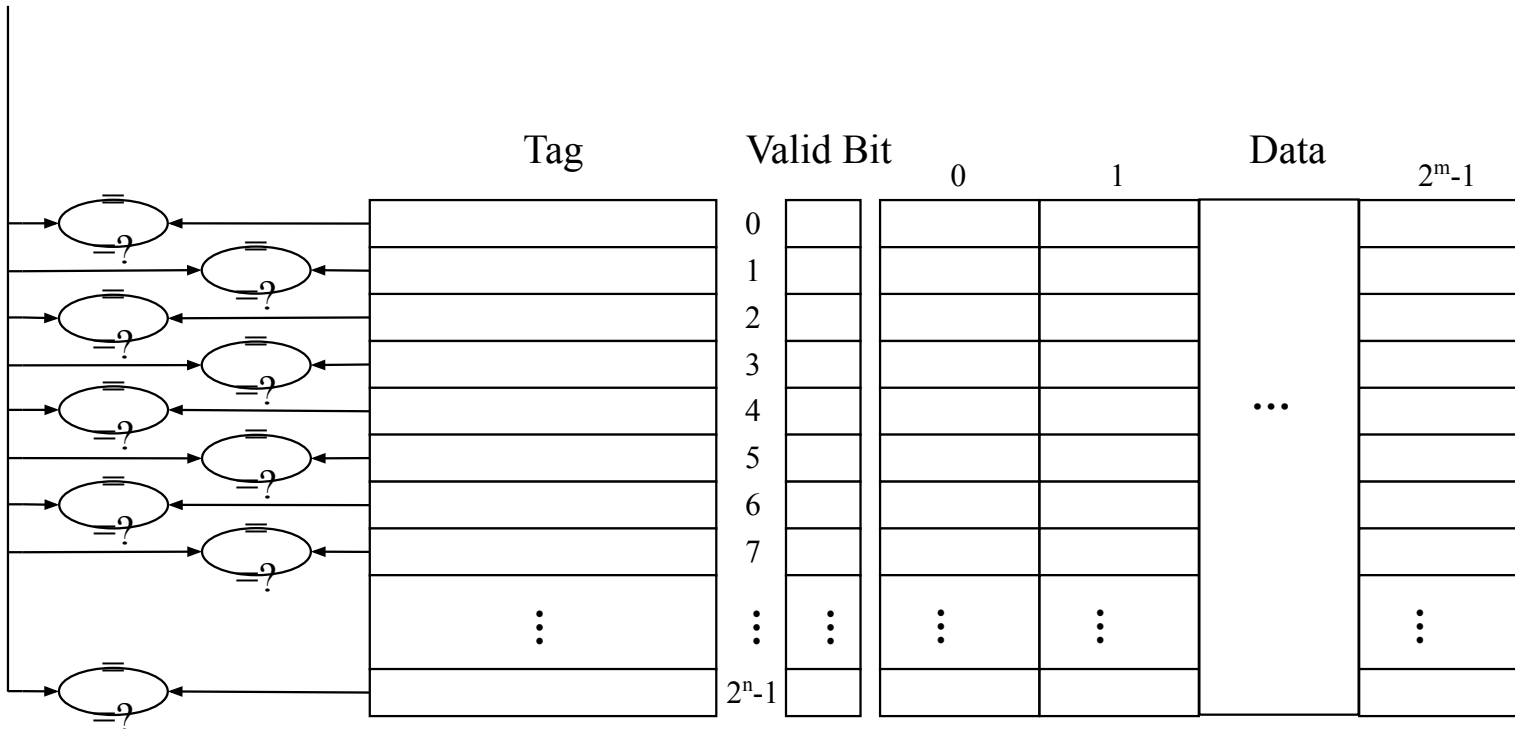


Associative Mapping

- Can handle up to N conflicts without eviction
- Better Temporal Locality
 - Assuming good Eviction policy
- More Complicated Control Logic
 - Slower to confirm hit or miss
 - Slower to find associated cache line

Fully Associative Cache

- Full Associative = 2^N Way Associative
 - Everything goes Anywhere!



Cache Arrangement

- **Direct Mapped** – Memory addresses map to particular location in that cache
 - 1-Way Associative Set
- **Fully Associative** – Data can be placed anywhere in the cache
 - 2^N -Way Associative Set
- **N-way Set Associative** – Data can be placed in a limited number of places in the cache depending upon the memory address

Replacement Methods

- If we need to load a new cache line, where does it go?
- Direct-mapped

Only one possible location

- Set Associative

N locations possible, optimize for temporal locality?

- Fully Associative

All locations possible, optimize for temporal locality?

What gets evicted?

- Approach #1: Random
 - Just arbitrarily pick from possible locations
- Approach #2: Least Recently Used (LRU)
 - Use temporal locality
 - Must track somehow – extra bits to recent usage
- In practice, Random ~12% worse than LRU

3 C's of Cache Misses

- Compulsory/Coldstart
 - First access to a block – basically unavoidable
 - For long-running programs this is a small fraction of misses
- Capacity
 - Had been loaded, evicted by too many other accesses
 - Can only be mitigated by cache size
- Conflict
 - Had been loaded, evicted by mapping conflict
 - Mitigated by associativity

Cache Miss Comparison

- Fill in the blanks: Zero, Low, Medium, High, Same for all

	Direct Mapped	N-Way Set Associative	Fully Associative
Cache Size: Small, Medium, Big?	Big (Few comparators)	Medium	Small (lots of comparators)
Compulsory Miss:			
Capacity Miss			
Conflict Miss			

Cache Miss Comparison

- Fill in the blanks: Zero, Low, Medium, High, Same for all

	Direct Mapped	N-Way Set Associative	Fully Associative
Cache Size: Small, Medium, Big?	Big (Few comparators)	Medium	Small (lots of comparators)
Compulsory Miss:	Same	Same	Same
Capacity Miss	Low	Medium	High
Conflict Miss	High	Medium	Zero

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

00001

00110

00001

11010

00110

Valid Bit

Tag

Data

0

1

2

3

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

000 01

001 10

000 01

110 10

001 10

0
1
2
3

Valid Bit

0
1
0
0

Tag

000

Data

M[00001]

Low N bits of address determine cache Line, rest is Tag (M=0 so no block address)

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

000 01

001 10

000 01

110 10

001 10

Valid Bit

Tag

Data

0
1
2
3

0
1
1
0

000
001

M[00001]
M[00110]

Access misses (compulsory), fills cache line 2

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

000 01

001 10

000 01

110 10

001 10

	Valid Bit	Tag	Data
0	0		
1	1	000	M[00001]
2	1	001	M[00110]
3	0		

Hits in cache

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

000 01

001 10

000 01

110 10

001 10

0
1
2
3

Valid Bit

0
1
1
0

Tag

000
001 != 110

Data

M[00001]
M[00110]

Miss (compulsory)

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

000 01

001 10

000 01

110 10

001 10

Valid Bit

Tag

Data

0

1

2

3

0
1
1
0

000
110

M[00001]
M[11010]

Evicts previous data from cache

Cache Access Example

Assume 4 byte direct-mapped cache (N=2, M=0)

Access pattern:

000 01

001 10

000 01

110 10

001 10

Valid Bit

Tag

Data

0

1

2

3

0
1
1
0

000
110!= 001

M[00001]
M[11010]

Miss (conflict) – this used to be in cache but was evicted

Cache Miss Example

- 8-word cache, 8-byte blocks. Determine types of misses (CAP, COLD, CONF).

Byte Addr	Block Addr	Direct Mapped	2-Way Assoc	Fully Assoc
0	0	Cold	Cold	Cold
4	0	Hit (in 0's block)	Hit (in 0's block)	Hit (in 0's block)
8	1	Cold	Cold	Cold
24	3	Cold	Cold	Cold
56	7	Cold	Cold	Cold
8	1	Conf (w/56)	Conf (w/56)	Hit
24	3	Hit	Conf (w/8)	Hit
16	2	Cold	Cold	Cold
0	0	Hit	Hit	Cap
	Total Miss:	6	7	6

Cache Miss Example

- 8-word cache, 8-byte blocks. Determine types of misses (CAP, COLD, CONF).

Byte Addr	Block Addr	Direct Mapped	2-Way Assoc	Fully Assoc
0				
4				
8				
24				
56				
8				
24				
16				
0				
	Total Miss:			

Cache Miss Example

- 8-word cache, 8-byte blocks. Determine types of misses (CAP, COLD, CONF).

Byte Addr	Block Addr	Direct Mapped	2-Way Assoc	Fully Assoc
0	0	Cold		
4	0	Hit (in 0's block)		
8	1	Cold		
24	3	Cold		
56	7	Cold		
8	1	Conf (w/56)		
24	3	Hit		
16	2	Cold		
0	0	Hit		
	Total Miss:	6		

With Remaining Time

- Finish Cache Example (Answers Below)
- Find out what your computer's caches are
- Write your own “Pop Quiz” on Caching
- We will exchange & take quizzes
 - Not graded
 - Then discuss

Quiz Ideas

- Describe a specific Cache:
 - How big is the tag? Where does address X map?
- Describe an access pattern:
 - Hit Rate? Average Access Time?

Cache Summary

- Software Developers must respect the cache!
 - Usually a good place to start optimizing
 - Especially on laptops/desktops/servers
 - Cache Line size? Page size?
- Cache design implies a usage pattern
 - Very good for instruction & stack
 - Not as good for Heap
 - Modern languages are very heapy

Matrix Multiplication

```
for (k = 0; k < n; k++){ for (i = 0; i < n; i++){  
    for (j = 0; j < n; j++){  
        c[k][i] = c[k][i] + a[k][j]*b[j][i];  
    }  
}}
```

Vs

```
for (k = 0; k < n; k++){ for (i = 0; i < n; i++){  
    for (j = 0; j < n; j++){  
        c[i][j] = c[i][j] + a[i][k]*b[k][j];  
    }  
}}
```

- Which has better locality?

Processor

Name	Intel Core i7 2600K		
Code Name	Sandy Bridge	Max TDP	95 W
Package	Socket 1155 LGA		
Technology	32 nm	Core Voltage	0.972 V



Specification Intel(R) Core(TM) i7-2600K CPU @ 3.40GHz

Family	6	Model	A	Stepping	7
Ext. Family	6	Ext. Model	2A	Revision	D2
Instructions	MMX, SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2, EM64T, VT-x, AES, AVX				

Clocks (Core #0)

Core Speed	1596.19 MHz
Multiplier	x 16.0 (16 - 38)
Bus Speed	99.76 MHz
Rated FSB	

Cache

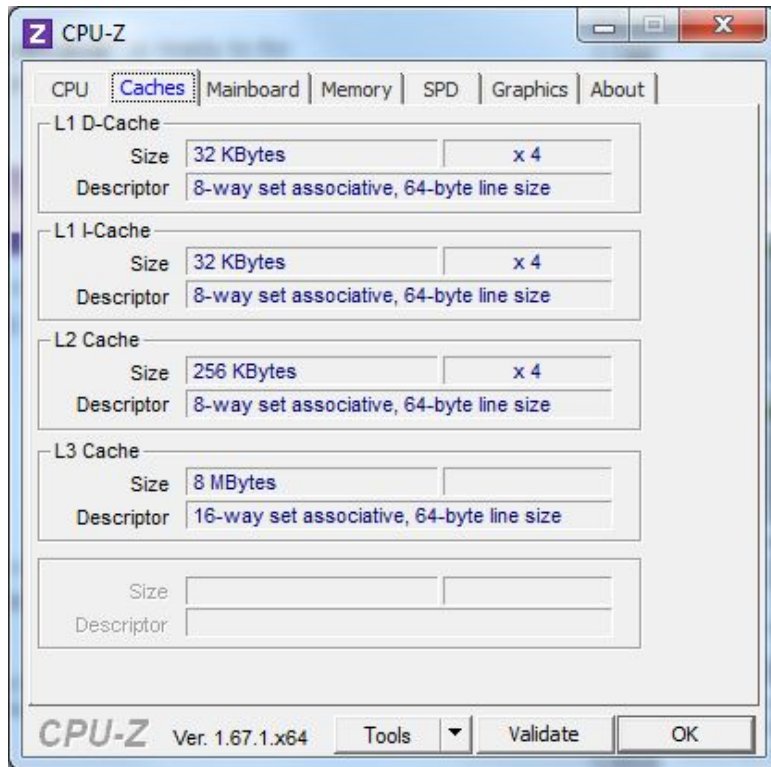
L1 Data	4 x 32 KBytes	8-way
L1 Inst.	4 x 32 KBytes	8-way
Level 2	4 x 256 KBytes	8-way
Level 3	8 MBytes	16-way

Selection Processor #1

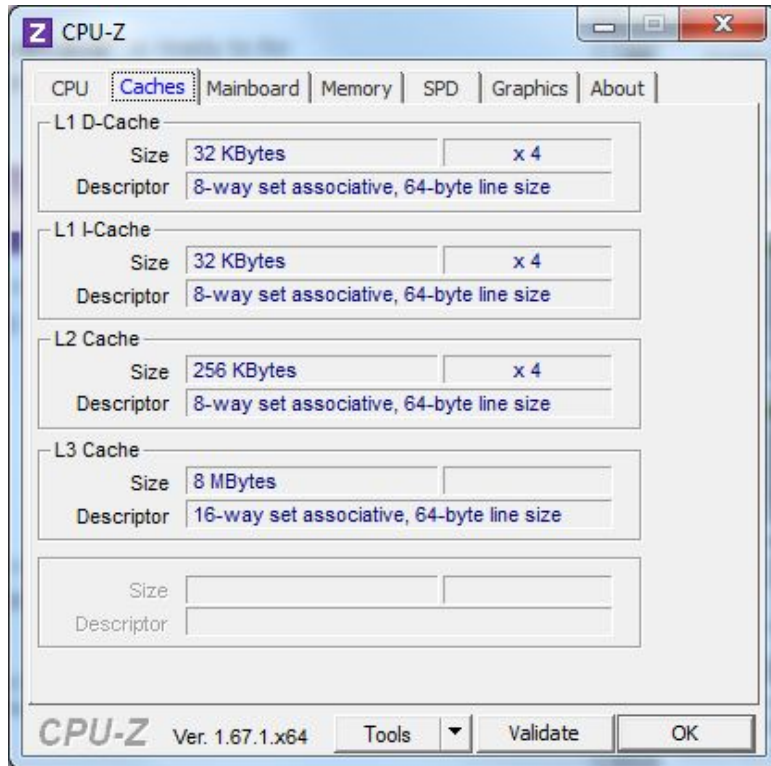
Cores 4 Threads 8

My L1 Data Cache

- $M = ?$
- $N = ?$
- Tag Size = ?



My L1 Data Cache



- M = 6 bits
 - 64-byte line size
 - 64 byte blocks
 - $2^6 = 64$
- N = 6 bits
 - 32kByte -> 15 bits
 - 8 way -> 3 fewer bits
 - M=6 -> 6 fewer bits
- Tag Size = 52 bits
 - 64 bit address space
 - Whatever is left over

Memory System Summary

Goal: provide the illusion of large, fast memory

Use hierarchy of memory types with complementary properties

Small/fast register file <-> huge/slow disk

Cache performance enabled by spatial and temporal locality in programs